



[Herrenkrugsteg Magdeburg]

# Advanced Topics in Feature-Model Analysis

## Thesis Topics and Software Projects

March 6, 2025

Elias Kuiter<sup>1</sup>

University of Magdeburg<sup>1</sup>

# Advanced Topics in Feature-Model Analysis

## 1. Introduction

## 2. Background

## 3. Thesis Topics

Extracting Feature Models From Extended KConfig Specifications (B/M)

CNF Transformation of Feature Models: A White-Box Analysis (B/M)

Topic Sketches

## 4. Software Projects

torte – Fully Automated Feature-Model Experiments in Bash (P)

clausy – Robust and Efficient CNF Transformation in Rust (P)

variED-NG – Collaborative Feature Modeling in TypeScript (P)

# 1. Introduction



Contact me:  
kuiter@ovgu.de

### About Me ...

- **until 2020:** M.Sc. Computer Science in Magdeburg
- **since 2021:** PhD student in Magdeburg supervised by Gunter Saake (Magdeburg) and Thomas Thüm (Braunschweig)

### Research Interests

- generally: software variability, quality, and evolution
- specifically: feature-model and product-line analysis
- i.e., software engineering  $\cap$  automated reasoning

### You ...

- are interested in working on cutting-edge research tools and topics
- successfully completed our **product-line course**

B: Bachelor thesis — M: Master thesis — P: software project

## 2. Background

# Modeling Features and their Dependencies

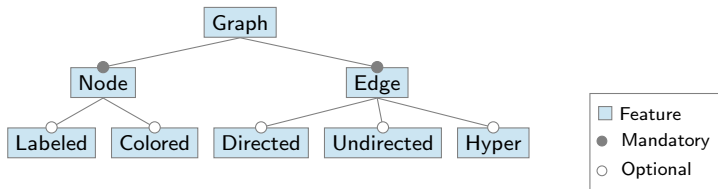
## Feature Models

- tree models **features**
- cross-tree **constraints** model dependencies
- solver-based **analyses** on the configuration space

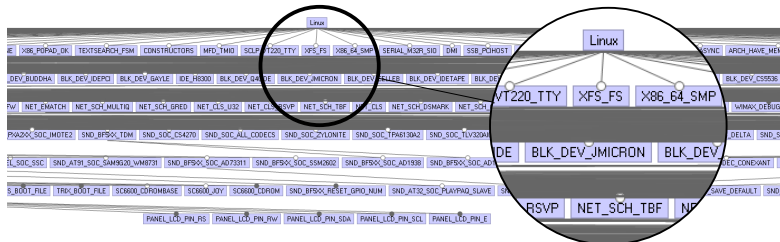
## The Linux Kernel

- $\approx 20,000$  features
- $> 10^{1,700}$  products
- 114 dead features [2013]
- 151 reverse dependency bugs [2019]

... and many other systems!



$$\neg(\text{Directed} \wedge \text{Undirected})$$
$$\text{Hyper} \rightarrow \text{Undirected}$$
$$\text{Directed} \nleftrightarrow (\text{Undirected} \wedge \text{Hyper})$$



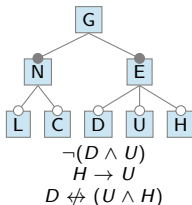
# Analyzing Feature Models with SAT and #SAT Solvers

## KConfig Specifications

```

config D
  bool "Directed"
  depends on !U
config U
  bool "Undirected"
  depends on !D
config H
  bool "Hyper"
  select U
    
```

## A Feature Model $FM$



## As a Formula $\Phi(FM)$

$G$   
 $\wedge (N \leftrightarrow G) \wedge (E \leftrightarrow G)$   
 $\wedge ((L \vee C) \rightarrow N)$   
 $\wedge ((D \vee U \vee H) \rightarrow E)$   
 $\wedge \neg(D \wedge U) \wedge (H \rightarrow U)$   
 $\wedge (D \not\leftrightarrow (U \wedge H))$

$\downarrow \Theta$

## Core Features

$\{G, N, E\}$

## Core Feature $F$ ?

$SAT(\Theta(\Phi(FM)) \wedge \neg F)$

## Feature Model Cardinality

8

## Products in $FM$ ?

$\#SAT(\Theta(\Phi(FM)))$

## As a CNF $\Theta(\Phi(FM))$

$\{\{G\}, \{\neg N, G\}, \{N, \neg G\},$   
 $\{\neg E, G\}, \{E, \neg G\}, \{\neg L, N\},$   
 $\{\neg C, N\}, \{\neg D, E\}, \{\neg U, E\},$   
 $\{\neg H, E\}, \{\neg D, \neg U\}, \{\neg H, U\},$   
 $\{\{D, U\}, \{D, H\}, \{\neg D, \neg U, \neg H\}\}\}$

# 3. Thesis Topics



# Extracting Feature Models From Extended KConfig Specifications (B/M)

## Problem

- currently, feature models cannot be extracted for every project that uses KConfig specifications
- i.e., some projects use the **Python-based parser KConfigLib** instead of a C implementation
- e.g., the feature model of the operating system **Zephyr** can currently not be analyzed

## Kconfig extensions

Zephyr uses the **Kconfiglib** implementation of **Kconfig**, which includes some Kconfig extensions:

- Default values can be applied to existing symbols without **weakening** the symbols dependencies through the use of `configdefault`.

## Goals

- review usage of KConfigLib in Zephyr and other projects in the system software domain
- investigate **differences between KConfigLib and the C parser** used in Linux
- conceptualize and implement a solution for **extracting feature models from KConfigLib** (e.g., transform into Linux-compatible KConfig, or adapt KClause to KConfigLib input)
- evaluate feasibility and performance + accuracy of extracted models by comparing with KClause

## Requirements

- reading literature and KConfig code
- experience with Python and C
- wrangle with subtle KConfig syntax/semantics

# CNF Transformation of Feature Models: A White-Box Analysis (B/M)

## Problem

- **CNF transformation** algorithms have a measurable impact on solver performance [Kuitert et al. 2022]
- there are many shades of CNF transformations, only few of which have been evaluated on a large corpus of models (**black-box analysis** 📦)
- a **white-box analysis** 🧰 is missing, which considers CNF parametrization and solver internals
- thus, we may better understand the influence of CNF and **choose good parameters**



## Goals

- extend our Rust tool `clausy` to with configurable CNF transformation (i.e., Tseitin threshold)
- instrument the C SAT solver `kissat` to monitor solver internals (i.e., pre-/inprocessing)
- evaluate performance of `clausy` + `kissat` for many parametrizations on a large corpus
- evaluate the minimum Tseitin threshold as a measure for feature-model complexity

## Requirements

- considerable experience with Rust and C
- evaluation design, reading literature

# Topic Sketches

On interest, we can develop these into concrete topics.

- cross-evaluate influence of CNF transformation algorithms implemented in different tools (B)
- evaluate the influence of CNF scrambling on solver performance (B)
- evaluate the influence of CNF transformation algorithms on sampling (B)
- extract solution-space models from KConfig, which encode program variants instead of configurations (M)
- evaluate construct validity of KConfig extractors (B/M)
- evaluate connection of CNF transformation with phase transition and community structure (B/M)
- ... (bring your own topic?)

## 4. Software Projects

# torte – Fully Automated Feature-Model Experiments in Bash (P)

## What is torte? 🍰

[@/ekuitier/🍰]

- a declarative workbench for **reproducible** feature-model analysis experiments
- can extract, transform, and analyze feature models in a **fully automated** fashion
- draft, execute distribute, and adapt experiments (without clone-and-own)

## A Simple Experiment: Counting BusyBox

```
experiment-subjects() {  
    add-busybox-kconfig-history --from 1_36_0 --to 1_36_1  
}  
experiment-stages() {  
    clone-systems  
    extract-kconfig-models  
    transform-models-into-dimacs  
    solve-model-count --timeout 10  
}
```

## Goal

fix problems and implement new features from roadmap (issue #1)  
⇒ enabling new use cases for torte

## Requirements

- experience with Bash programming
- some experience with Docker
- willing to write clean code in Bash :-)



# clausy – Robust and Efficient CNF Transformation in Rust (P)

## What is clausy? 🧑🏻

[[@ekuiter/](#)🧑🏻]

- transforms feature-model formulas into **conjunctive normal form (CNF)** for subsequent analysis
- improves comparability of transformation algorithms, competes with state-of-the-art tools
- implements algorithms for advanced use cases

## Some Examples: Classify Feature-Model Formulas

# simplify a given CNF

```
bin/clausy model.dimacs
```

# read from standard input and count solutions

```
cat model.uvl | bin/clausy --.uvl to_cnf_tseitin count
```

# prove model equivalence

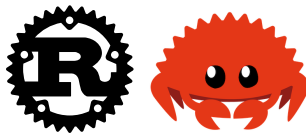
```
! bin/clausy a.model b.model '+(*(-1 2) *(1 -2))'  
to_cnf_tseitin satisfy &>/dev/null
```

## Goal

fix problems and implement new features from roadmap (issue #1)  
⇒ enabling new use cases for clausy

## Requirements

- programming experience in Rust
- interested in algorithm engineering



# variED-NG – Collaborative Feature Modeling in TypeScript (P)

## What is variED?

[[ekuitier/variED](https://github.com/ekuitier/variED)]

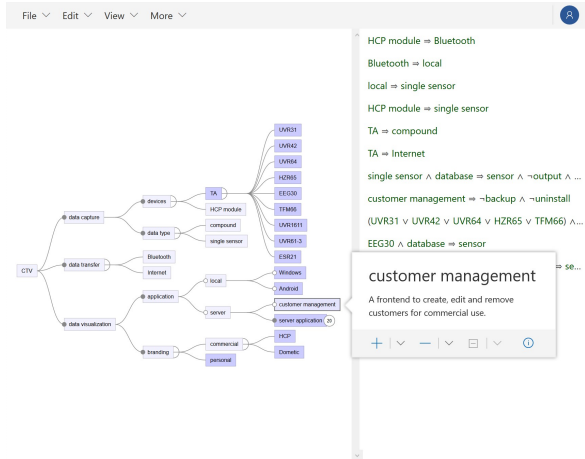
- a research prototype for viewing and editing feature models in a **collaborative, real-time** fashion
- vision: “Google Docs” for feature modeling
- easily share models, teach modeling, ...

## Goal

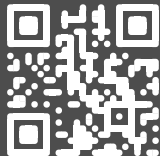
- prototype is complex and currently defunct
- continue work on simplified fork variED-NG
- fix and improve the editor

## Requirements



- experience with TypeScript + React/Redux
- interested in modern web development





Interested?



Contact me: [kuiter@ovgu.de](mailto:kuiter@ovgu.de)

/ekuiter/

/ekuiter/



## 5. Archived Topics

These topics have already been assigned to or completed by another student.  
They can give you an impression of what we have done in the past.  
Some of these topics might also be suited for follow-up work.

# The Impact of Clone-and-Own on Parsing KConfig Specifications (B/M)

## Problem

- the C KConfig parser of Linux **evolves over time**, slowly adding new syntax and semantics and deprecating or fixing old constructs
- moreover, other projects tend to **clone-and-own** the parser from Linux, leading to a variety of implementations
- it is unknown which (distinct) KConfig implementations exist and how they came to be
- it is also unclear to which degree these implementations are **compatible** with each other, and whether using an incorrect parser can threaten the validity of research evaluations

## Goals

- trace the **genealogy of KConfig parsers** (who forked it when/why? with which adaptations?)
- evaluate **backwards compatibility** of the parser in a given project (can a new parser extract an old feature model? is the result accurate?)
- evaluate **cross-compatibility** of parsers

## Requirements

- reading literature, code, and commit histories
- experience with C to compare implementations
- wrangle with subtle KConfig syntax/semantics



# Minimizing CNFs to Isolate Solver Bugs (B)

## Problem

- CNFs of real-world feature models sometimes uncover **bugs** even in production-grade (#)SAT and SMT solvers
- e.g., in countAntom, sharpSAT/dSharp, Z3, clausy, FeatJAR
- during development and maintenance of such solvers, reducing problematic CNFs to a **minimum non-working example** can facilitate finding the causes of bugs, reporting them, and preventing future regressions
- however, this process is currently a manual task and time-consuming



## Goals

- identify fault oracles (e.g., solver crashes), re-view reduction strategies (e.g., removing clauses one-by-one, bisection, backtracking to avoid a local minimum)
- implement a (semi-)automatic tool that repeatedly reduces clauses and literals in a faulty CNF until it is minimal
- evaluate performance and compare with global minimum (e.g., obtained manually)

## Requirements

- algorithm design, reading literature
- potential challenges: generative effects, local minima, choosing bugs to evaluate

# A Dashboard for Evolving Variability in Open-Source Systems (P)

## Problem

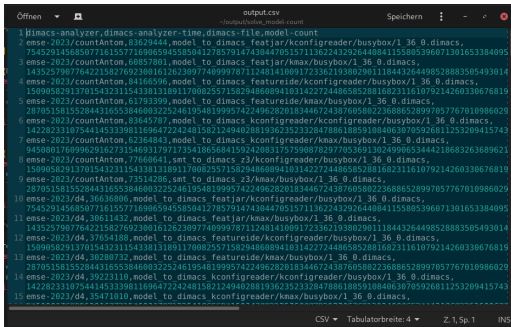
- torte fully automates feature-model analysis
- can be used to analyze latest Linux kernel
- **but**: no user-friendly frontend exists yet

## Goal

- develop a web frontend for torte
  - find appropriate visualizations
- ⇒ quick visualization of current state of variability  
⇒ extension of previous research evaluation

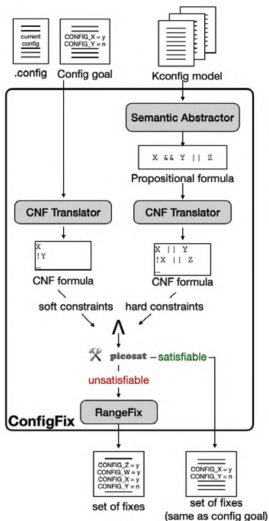
## Requirements

- experience with frontend development (e.g., HTML/CSS, React/Vue/Dash, ...)
- backend experience not necessarily needed (e.g., assume a static CSV file over AJAX)



```
Öffnen  output.csv  Speichern
-./output/sove_model-count
1|dimacs-analyzer,dimacs-analyzer-time,dimacs-file,model-count
2|emse-2023/countAntom,83629444,model to dimacs featjar/kconfigreader/busybox/1 36 0,dimacs,
754529145685077161557716906594558504127857914743044705157113622432926448841155805396071301653384095
3|emse-2023/countAntom,60857801,model to dimacs featjar/kmax/busybox/1 36 0,dimacs,
143525796776422158276923061612623097740999787112481410091723362193802901118443264498528883505493014
4|emse-2023/countAntom,84166596,model to dimacs featureide/kconfigreader/busybox/1 36 0,dimacs,
150905829137015432311543381318911700825571582948608941031422724486585288168231161079214260330676819
5|emse-2023/countAntom,61793399,model to dimacs featureide/kmax/busybox/1 36 0,dimacs,
28705158155284431655384600325246195481999574224962820183446724387605802236886528997057767010986029
6|emse-2023/countAntom,83645787,model to dimacs kconfigreader/kconfigreader/busybox/1 36 0,dimacs,
14228233107544145333981169647224248158212494028819362352328478861885910840638705926811253209415743
7|emse-2023/countAntom,62364843,model to dimacs kconfigreader/kmax/busybox/1 36 0,dimacs,
94508017609962916273154693179173541865684159242083175759087829770536913024990653444218683263689621
8|emse-2023/countAntom,77660641,smt to dimacs z3/kconfigreader/busybox/1 36 0,dimacs,
150905829137015432311543381318911700825571582948608941031422724486585288168231161079214260330676819
9|emse-2023/countAntom,73514286,smt to dimacs z3/kmax/busybox/1 36 0,dimacs,
28705158155284431655384600325246195481999574224962820183446724387605802236886528997057767010986029
10|emse-2023/d4,36636806,model to dimacs featjar/kconfigreader/busybox/1 36 0,dimacs,
754529145685077161557716906594558504127857914743044705157113622432926448841155805396071301653384095
11|emse-2023/d4,36611432,model to dimacs featjar/kmax/busybox/1 36 0,dimacs,
143525796776422158276923061612623097740999787112481410091723362193802901118443264498528883505493014
12|emse-2023/d4,37654188,model to dimacs featureide/kconfigreader/busybox/1 36 0,dimacs,
150905829137015432311543381318911700825571582948608941031422724486585288168231161079214260330676819
13|emse-2023/d4,38280732,model to dimacs featureide/kmax/busybox/1 36 0,dimacs,
28705158155284431655384600325246195481999574224962820183446724387605802236886528997057767010986029
14|emse-2023/d4,39223110,model to dimacs kconfigreader/kconfigreader/busybox/1 36 0,dimacs,
14228233107544145333981169647224248158212494028819362352328478861885910840638705926811253209415743
15|emse-2023/d4,35471010,model to dimacs kconfigreader/kmax/busybox/1 36 0,dimacs,
```

# Evaluating an Extractor for KConfig-Based Feature Models (B)



## Problem

- a new feature-model extractor, ConfigFix, has recently been published
- however, it has not been compared to existing extractors yet

## Goal

- implement extraction of feature models with ConfigFix in torte
- evaluate efficiency and accuracy on a large corpus on feature models

## Requirements

- good in bash programming
- work with ConfigFix, which is implemented in C

# Extracting Feature Hierarchies for KConfig-Based Feature Models (B)

## Problem

- feature-model extractors for KConfig mostly ignore the **feature hierarchy**
- tooling for extracting hierarchies is now defunct, identification of **feature parents** in Kconfig is yet under-researched

```
1 namespace Root
2
3 features
4     Root
5         optional
6             UNWINDER_ORC
7             UNWINDER_FRAME_POINTER
8             UNWINDER_GUESS
9             X86_64
10            IO_DELAY_0X80
11            IO_DELAY_0XED
12            IO_DELAY_UDELAY
13            IO_DELAY_NONE
14            BRANCH_PROFILE_NONE
15            PROFILE_ANNOTATED_BRANCHES
```

## Goal

- extract a feature hierarchy from KConfig specifications + evaluate accuracy
- and/or: reverse-engineer hierarchy from formula + compare with KConfig hierarchy

## Requirements

- adjusting KConfig parser written in C
- adjust or implement a tool for reverse-engineering
- c.f. [Yaman 2023](#), [Yaman et al. 2024](#)

# Feature-Model Analysis with SAT Solvers: A Journey Through Time (B)

## Problem

- feature models grow more complex over time
- automated reasoning tools (e.g., SAT solvers) get more efficient over time
- **but**: which development is faster? can SAT solvers actually keep up?

## Goal

- collect best SAT solvers of the last 20 years
- collect feature models from the last 20 years
- run selected feature-model analyses with solver from year **X** on model of year **X**
- evaluate evolution of SAT solving performance (cf. Moore's law)
- see time leap challenge

## Requirements

- methodology design, reading literature
- challenges: data availability and formats



[Photo: Laurent Simon]



# Previously Supervised Theses [full list available [here](#)]

- A Non-Clausal Slicing Algorithm for Feature Models (M)
- Efficiency of Projected Model Counting for Feature-Model Analysis (B)

Excerpts of feature models (so-called slices) allow for modularization and separation of concerns for software variability. However, calculating slices is complex. For some applications, it is sufficient to know the size (i.e., the number of valid products) of a slice. Projected Model Counting (PMC) is a new solver class that can directly determine the size of a slice without calculating it. The aim of this work is to evaluate the applicability of PMC for feature model analysis. Tasks include to research PMCs and their functionality, apply PMC to feature models of various sizes, and compare with slicing and subsequent #SAT calls.

- Evaluating the Efficiency of Hybrid CNF Transformations for Feature-Model Formulas (M)

For automated feature-model analysis, SAT and #SAT solvers are used. These expect the input formula to be in CNF (conjunctive normal form). There are various methods to convert a formula into CNF. The aim of this work is to continue our existing research and conduct a detailed comparison of the distributive transformation with the Tseitin and Plaisted-Greenbaum transformations, thereby drawing conclusions about when to prefer which CNF transformation. Some aspects to consider are parameters for hybrid transformation, preprocessing in SAT solvers, and multiple implementations for external validity.

- Configuration of Software Product Lines With Configuring Constraints via Feature Attributes (M)
- A Semi-Automated Release Management Process for Microservices (B)
- Reengineering of a Microservice Architecture: A Case Study on the PEGASOS System (M)
- System Tools' Firmware Flashing Automation in a Complex Software/Hardware Environment (M)